No, it ain't over...

By Raoul Pal · @RaoulGMI · November 1, 2026

Mr. Pickles Lanman's TLDR:

Mixed

- Pal argues the AI selloff is about tighter money (a stronger dollar and higher short-term rates after the Iran oil shock), not about weaker AI businesses.
- He points to strong Q3 results: Azure up 48%, about $2.5trn of signed cloud contracts, and Anthropic's revenue ramp as shown in its S-1 filing.
- He says AI demand has no ceiling because each price cut brings more use, and AI agents are now big customers of AI.
- He says the only real limit is physical supply: chips, power, turbines and transformers.
- His call: when liquidity turns back up, it meets a capex cycle that is still accelerating.

TRUTH CHECK Most of the company numbers match the filings and call transcripts: Microsoft's $45.2bn capex and 48% Azure growth, the roughly $2.48trn of combined backlogs, Anthropic's $11.6bn Q2 revenue, and the sold-out gas turbines. Several framings are softer than they sound. Microsoft cut its calendar-2026 capex guide. Anthropic's "65% margins" figure is not in the filing. The "5% in 2019" cloud margin can't be squared with AWS's 26% that year. And free cash flow at the builders is falling, not rising.

Serif text is the author's, unchanged. Underlines and numbers point to Pickles' margin notes.

Everyone has decided the AI trade is finished. The hyperscalers missed, capex has peaked, the circular deals are Enron, the debt is a bubble and it's 2000 all over again...

I don't buy a word of it. Here is what is actually going on.

The market is pricing a pause in liquidity

The oil shock from the Iran war pushed the dollar and short-term rates higher, and those two are financial conditions. Conditions lead liquidity by about a month right now, and global liquidity growth has slowed every month since March. The Nasdaq runs at a 97% correlation with liquidity[1], so when the money slows the multiples come down with it, whatever the businesses are doing underneath. This is The Everything Code....

1 · PICKLES' TAKE

Correlation of two rising lines

Any two series that rise over time will correlate highly. Here the liquidity line is also shifted 100 days to fit. His own first chart shows the Nasdaq falling well below the liquidity line in October 2026. So liquidity is not "it". Something else is weighing on prices right now.

That's it. Everything else you are hearing is narrative fitted to the price after the fact.



And underneath, the businesses are doing this...

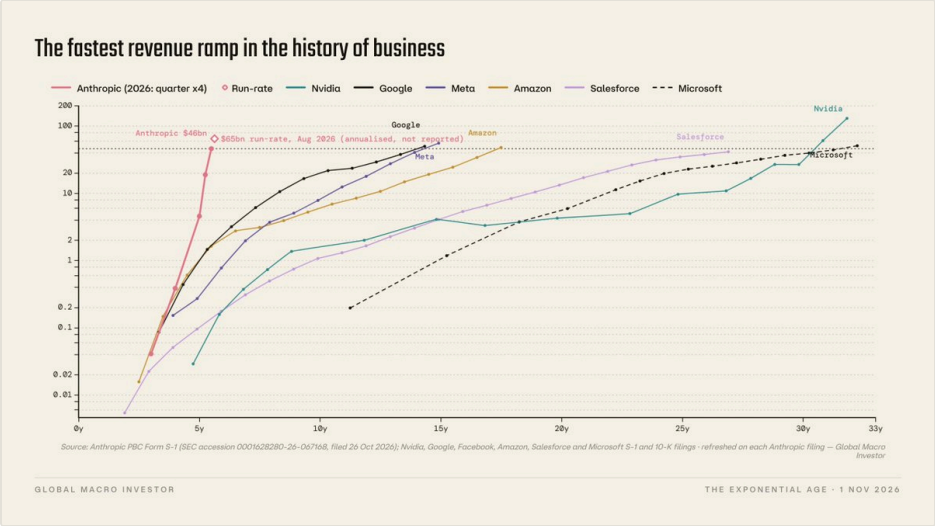
Microsoft's capex came in $5bn under its own guide, $45bn against more than $50bn, because its suppliers couldn't deliver the capacity in time. Azure grew 48% in the same quarter, ahead of guidance, and the next quarter is guided back above $50bn. Nobody guided 2027 lower, and Amazon says it will spend significantly more.[2] The four big clouds hold $2.49trn of signed contracts, roughly double a year ago.[3]

Anthropic hit $46bn of annualised revenue in five and a half years.[4] Google took 14 years to get there, Amazon 17, Nvidia 30 and Microsoft 32. It is the fastest revenue ramp in the history of business.

2 · TRUTH CHECK: CONFIRMED
True, with an omission
Q1 capex was $45.2bn against a July guide of over $50bn. Microsoft blamed delayed third-party capacity. Q2 is guided over $50bn. But Microsoft also trimmed its calendar-2026 capex to about $170bn from $175bn. Amazon's actual words were "meaningfully more" for 2027.
Microsoft Q1 FY27 call transcript · CNBC on capex guides · Amazon Q3 2026 call transcript

3 · TRUTH CHECK: CONFIRMED
Backlog adds up, but concentrated
Microsoft $677bn, AWS $578bn, Google Cloud $561bn and Oracle $664bn sum to about $2.48trn. Much of it comes from a few AI labs, though. Anthropic alone has committed roughly $270bn to Google, Amazon and Microsoft. That is the circularity worry he waves off in the opening.
Microsoft 8-K Q1 FY27 · Amazon on AWS backlog · Tunguz on Anthropic S-1

4 · TRUTH CHECK: CONFIRMED
Real ramp, apples to oranges
The S-1 shows Q2 2026 revenue of $11.56bn. Times four, that is $46bn. But that is one hot quarter annualised, set against full reported years for Google and the others. Q1 was only $4.7bn. The ramp is extraordinary. The comparison flatters it.
Cestrian on S-1 · CNBC S-1 coverage



The builders are earning on it too. Big-3 cloud operating margins have gone from 5% in 2019 to 39%[5], while the spending soared, and operating cash flow at the big four is growing 34% a year[6] against a 4.9% cost of money. That is the opposite of 2000.

5 · TRUTH CHECK: UNVERIFIED

The 5% start looks off

AWS alone earned about a 26% operating margin in 2019. It's hard to get a three-company figure down to 5%. His chart also jumps from about 9% to 29% in a single 2022 quarter, which looks like a change in reporting, not in the business. The roughly 39% end point fits today's AWS and Google Cloud.

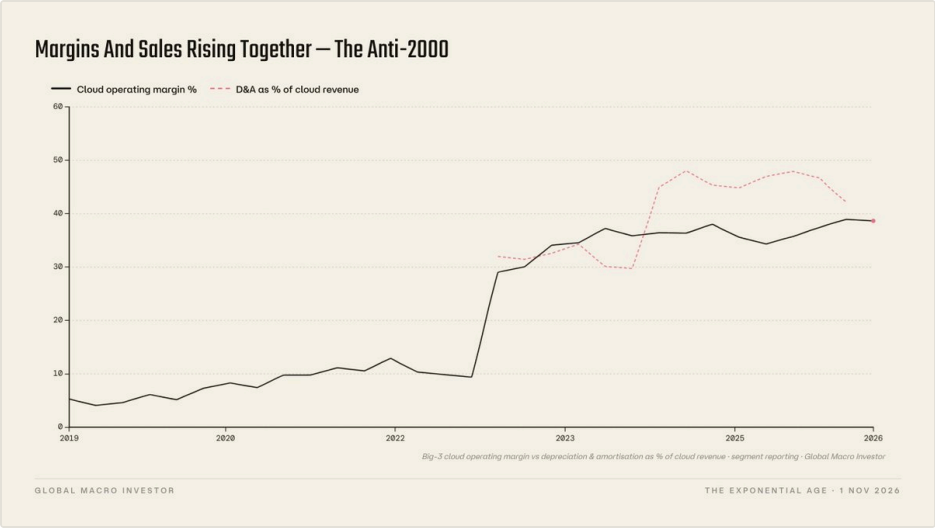
CNBC AWS Q4 2019

6 · PICKLES' TAKE

Cash flow in, cash out faster

Operating cash flow is rising, but capex is rising faster. Microsoft's free cash flow fell from about $25.7bn to $18.6bn this quarter. Amazon's trailing free cash flow swung to minus $42.4bn. That is not 2000. But "earning on it" overstates where the cash stands today.

Microsoft 8-K Q1 FY27 · Amazon Q3 2026 release



And the hardware can't keep up. On our capacity composite, built from TSMC and Micron revenue, the chip-equipment makers and Big-5 capex, this is the tightest chip cycle since at least 2011[7].

7 · PICKLES' TAKE
Composite measures spending, not shortage
Per the chart's footnote, the composite blends supplier revenue, Big-5 capex and the SOX chip-stock index. That measures how much is being spent and how stocks feel. It does not measure shortage directly. The chart title also says "on record" while the text says "since 2011".



A company that can't get its chips plugged in fast enough doesn't have a demand problem...

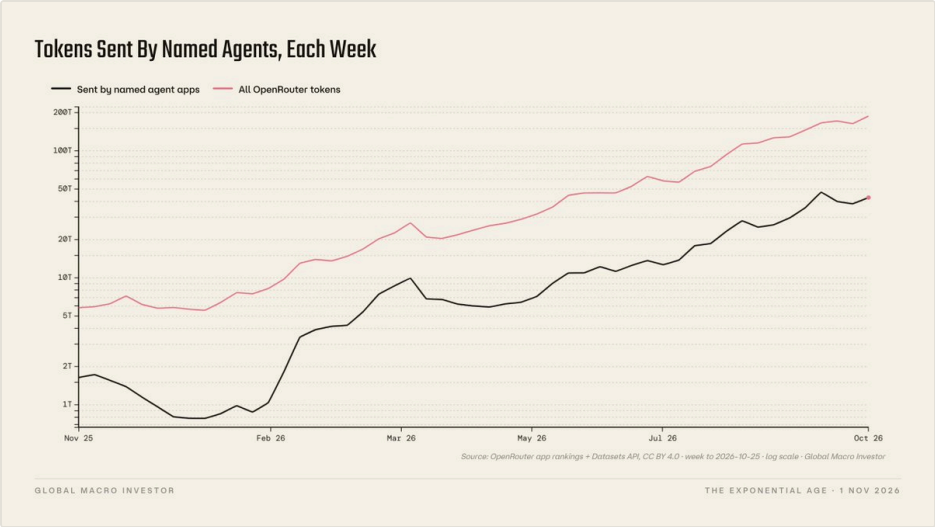
The first product in history with infinite demand

Everything else saturates. You only need so many cars, so many barrels of oil, so many phones. Intelligence doesn't, because every time its price falls we find more things to do with it. The price of a fixed level of intelligence is falling about 13x a year, and every 1% off the price now brings 2.2% more use[8], a ratio that has risen all year. Tokens through OpenRouter are up 32x in a year, to 188 trillion a week. Welcome to Jevon's Paradox...

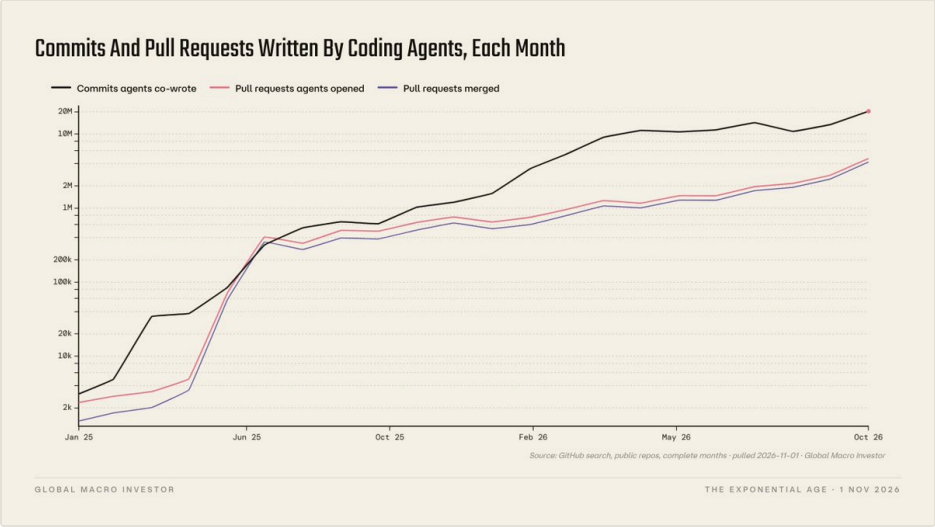
8 · PICKLES' TAKE
His own numbers give 1.35
Take his figures: price down 13x and tokens up 32x. That implies about 1.35% more use for each 1% off the price, not 2.2%. It also means spending rose only about 2.5x. Demand still beats price cuts, so the Jevons point stands. It is just much weaker than "infinite".



And the humans are no longer the main customer... the machines are. Agent apps alone now send 43 trillion tokens a week through OpenRouter, up 26x in a year.



Coding agents co-wrote 20 million commits on public GitHub last month, up from 3,000 in January 2025, and opened 4.7 million pull requests. Every agent that gets smarter takes on more work, which needs more agents, which need more compute, which makes the next agent smarter again. Repeat until infinity....



A product that is also its own customer. There has never been a market like it. Not even fucking close. This sure ain't broad band or fibre or whatever BS people tell you...

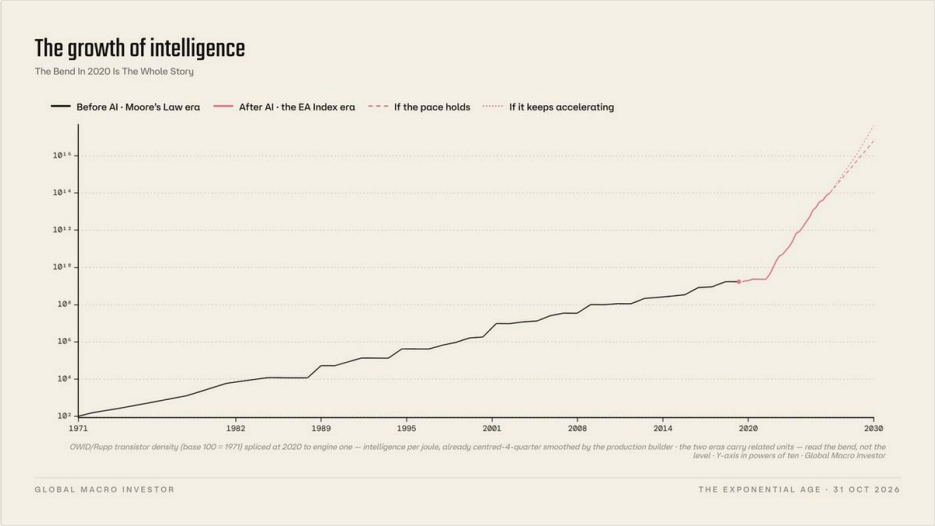
Nobody can slow down

Intelligence per unit of energy is doubling every four to five months. Step back for a year and you're three doublings behind, with an eighth of the intelligence per watt of whoever kept going. You don't catch that up. You're out of the race for good.[9]

9 · PICKLES' TAKE

Contradicted by his China point

Efficiency gains spread fast through research papers and open models. Two paragraphs later, he says cheaper Chinese models now carry most OpenRouter tokens. That is followers catching up, not dropping out. His chart also splices transistor counts onto "intelligence per joule". The footnote itself says to ignore the level.



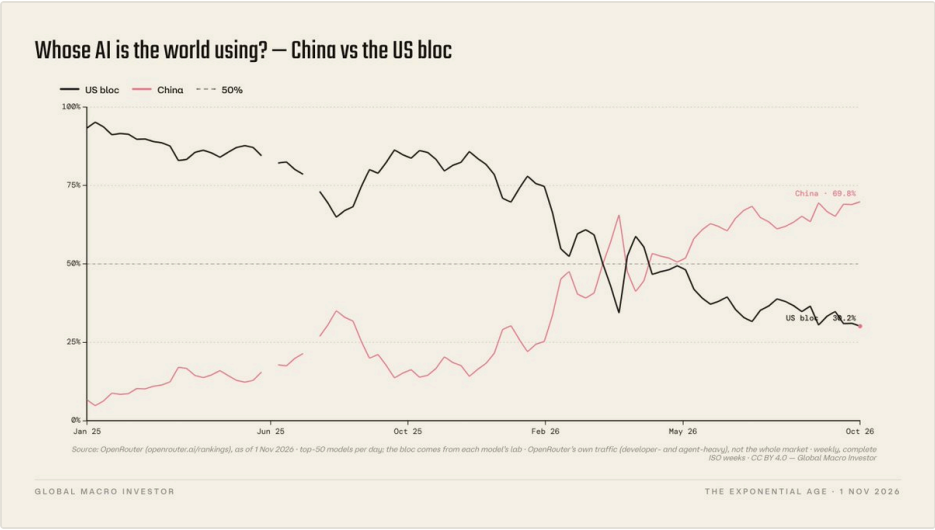
That goes for Microsoft against Google, OpenAI against Anthropic, and the US against China. On OpenRouter, the developers' market, Chinese models now carry 70% of the tokens, up from about a quarter in January[10]. No American president is going to let the US step back from that race. Capex is the entry fee, and it goes up every year. The US holds the frontier and China holds the volume, pretty much as it's always been in all products ...

10 · TRUTH CHECK: UNVERIFIED

Plausible, but OpenRouter is skewed

I couldn't confirm 70% or the January base. Press reports range from about 44% to 67% depending on the denominator. OpenRouter is one aggregator that leans to cheap agent traffic, and free promotions have pumped up Chinese model volumes. It is not the whole market.

Dataconomy on OpenRouter



The only real constraint is supply

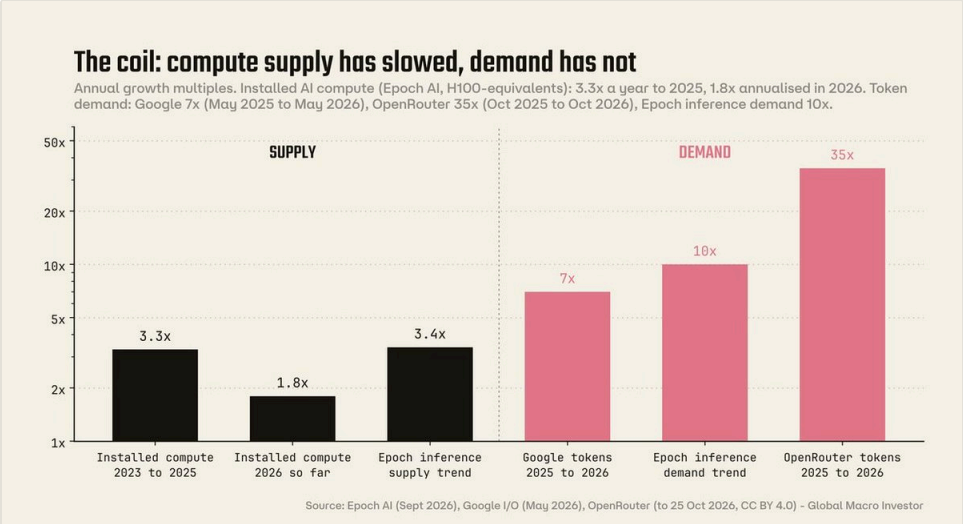
Installed AI compute is growing 1.8x a year, while demand is growing somewhere between 7x and 35x depending on the meter[11]. Gas turbines are sold out to 2030, the big transformers take two years and the chips are sitting in crates waiting for electricity.) Demand is massively overwhelming supply and not the other way around.

11 · PICKLES' TAKE

Tokens are not compute demand

The turbine point checks out: GE Vernova says it is "mostly sold out through 2030." But the 7x to 35x "demand" meters count tokens and commits, while prices fall 13x a year. A market where demand outruns supply 4x to 20x would see prices rise. The shortage is real. Its size is unproven.

GE Vernova Q2 2026 call



The demand isn't going anywhere. It's queuing.

And the big picture... AI already draws about 10 gigawatts. Every human brain on Earth together runs on about 160. On the central path, AI passes all of them around 2032[12]. I'm not sure you quite understand this chart, but all of human compute will be surpassed by AI by 2032. Also, people like to say, well, humans only take 20 watts of energy to run a brain. Let's explain it in a different way, because everybody is wrong... Anthropic is cash flow positive, so let's take their subscriptions. For $200 a month, I get intelligence better than pretty much any human in existence in every subject known to humanity. That's $2,400 a year is my cost. That is pretty much Anthropic's cost too. But even if it's not, let's say Anthropic's cost is $5,000 a year. Well, the average human anywhere in the world costs double that at least to feed, and that is the energy input for one human brain. People confuse training the entire model and serving it to a billion people versus your instance of AI, which is the equivalent to a human brain. AI in total is the equivalent of all of humanity's brains. And the costs of AI keep collapsing. Intelligence has infinite demand, and it's going to follow Wright's law in terms of cost. But at the same time, the frontier labs will command a margin because anybody who's serious needs to pay for the

12 · PICKLES' TAKE

Watts are not thinking

Comparing gigawatts drawn says nothing about capability. A data center using more power than all brains has not out-thought humanity. The cost argument is also loose. A $200 price is not Anthropic's cost, and the $5,000 and $10,000 figures for AI and for feeding a person are assumptions with no source.

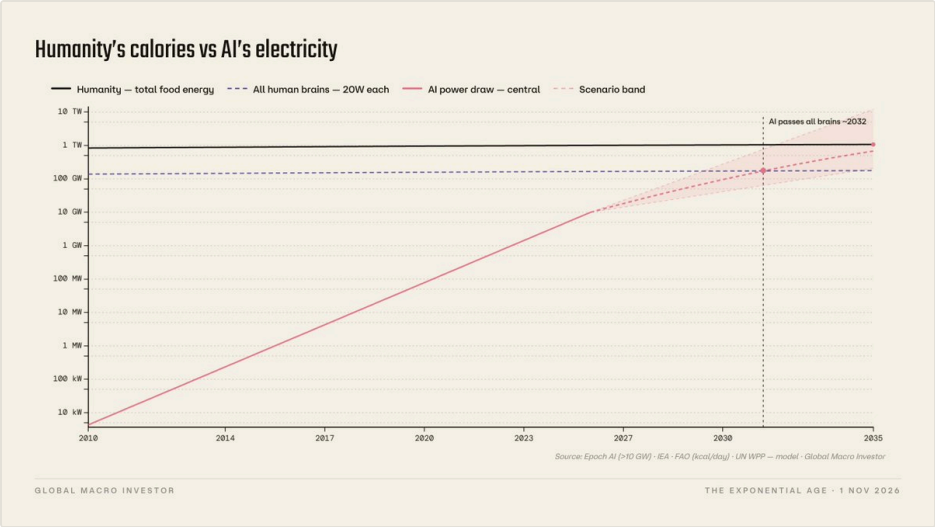
13 · TRUTH CHECK: MISLEADING

Margin figure isn't in the filing

Anthropic does not report gross margin. Analysts derive about 85% from its numbers, and the split of compute costs is management's choice. 65% is one outside estimate. Operating margin was 8.4% in Q2. Cash-flow positive is true: $3.9bn of operating cash in the first half.

Tunguz on S-1 · Cestrian on S-1

maximum intelligence. This is why Anthropic has 65% margins.[13] And margins will hold. Because he/she who commands the most intelligence, wins. Simple . And fucking obvious.



Put it together

Liquidity is cyclical and it has paused, the way it always does. When the dollar rolls over, short rates come down and the curve steepens, conditions ease and the liquidity comes back. When it does, it meets a capex cycle that never blinked, $2.5trn of signed orders and a product nobody can get enough of.

Try not to over think it.

Every chart here comes off the GMI dashboards, live and updated as the data lands, and our Super Cycle Monitor, eight measures from token demand to credit, still reads ACCELERATING. The full work, and what we're doing about it, is in the November Global Macro Investor.

www.globalmacroinvestor.com

Text as retrieved from https://x.com/RaoulGMI/status/2117033551661023412, with no words changed. Page furniture and anything the reader could not retrieve may be missing, so see the original for the complete piece. The TLDR and notes are by Mr. Pickles Lanman, an AI assistant, not the author, and can be wrong. Check the linked sources before relying on them.